


Deep Learning

MSiA 490-30



Theory and Applications



NORTHWESTERN
UNIVERSITY



Preliminaries

- Class photo?
- GPU/AWS?
- Heatmaps
- Groups
- Projects



NORTHWESTERN
UNIVERSITY

MSIA 490-30: Deep Learning. Spring 2017.
Instructor: Dr. Ellick Chan. TA: Mark Harmon.

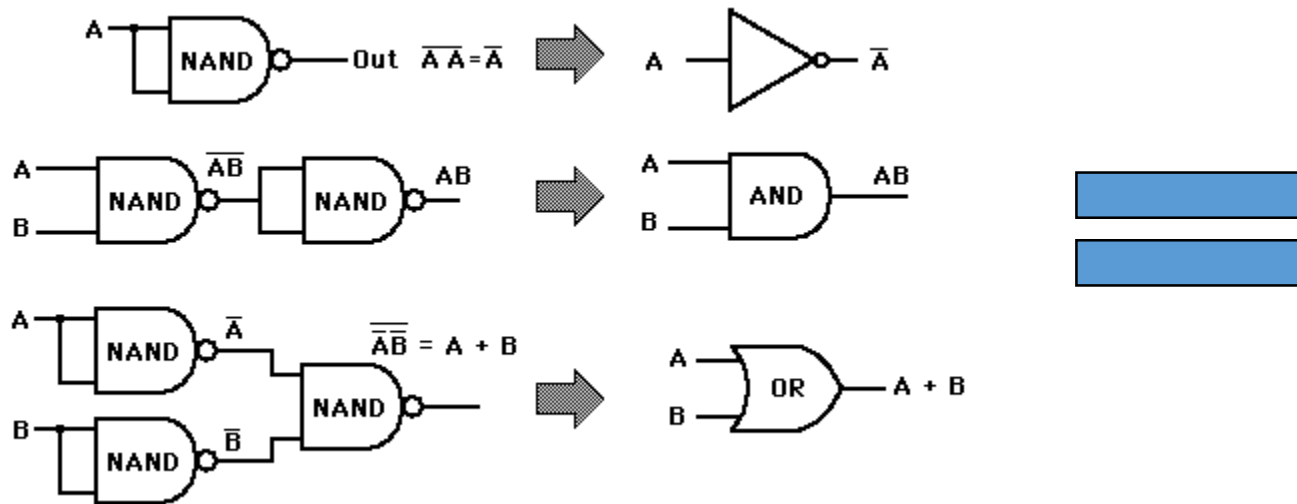
Feedback: News and Ideas



~5-10 min to share ideas/news



NAND gates universal for computation



Most modern
computers,
phones, tablets



Last time...


- Python tutorial
 - Linear classifier identifying 3's
- This time:
- How to improve classification performance
 - NN in “depth”






NORTHWESTERN
UNIVERSITY

History: Perceptrons



Schematic of a biological neuron.



Schematic of Rosenblatt's perceptron.




Rosenblatt, 1957

Think of NN as statistical
generalization machines with
priors - not as brain models



NORTHWESTERN
UNIVERSITY

“Layered” Logistic Regression






Recall...

- Simple Linear classifier
- Creates a separating hyperplane defined by:

$$f_w(x) = w_0 + w_1x_1 + \cdots + w_px_p$$





Improving predictions

- To “improve” predictions, first define an objective function
- Any ideas?
- Hint: Reward for correct answers, penalize for wrong



Improving predictions

- To “improve” predictions, first define an objective function
- Simple penalty: Mean Square Error (MSE)

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - t_i)^2$$

over n examples

Correct Label (t)	Prediction (y)	Penalty
0	0	0
0	1	1
1	0	1
1	1	0



Improving predictions

- $MSE = \frac{1}{n} \sum_{i=1}^n (y_i - t_i)^2$
- Recall simple linear classifier:
$$f_w(x) = 1 * w_0 + w_1 x_1 + \dots + w_p x_p$$
$$= \sum_{j=0}^p w_j x_j = w \cdot x$$
- $MSE = \frac{1}{n} \sum_{i=1}^n (f_w(x_i) - t_i)^2$
- $MSE = \frac{1}{n} \sum_{i=1}^n \left(\sum_{j=0}^p w_j x_j - t_i \right)^2$

Any ideas how to minimize MSE?



Improving predictions

- $MSE = \frac{1}{n} \sum_{i=1}^n (y_i - t_i)^2$

- Recall simple linear classifier:

$$f_w(x) = 1 \leftarrow w_0 + w_1 x_1 + \dots + w_p x_p$$
$$= \sum_{j=0}^p w_j x_j = w \cdot x$$

Remember that
 $x_0=1$

This is called the
“dot product”

- $MSE = \frac{1}{n} \sum_{i=1}^n (f_w(x_i) - t_i)^2$

- $MSE = \frac{1}{n} \sum_{i=1}^n \left(\sum_{j=0}^p w_j x_j - t_i \right)^2$

Any ideas how to minimize MSE?




Simple example

- $y = f(w, x)$
 $= w_0 + w_1x_1 + w_2x_2$

x	t	$y=f(x)$
$\sim(3,2)$	1	~ 1
$\sim(-2,-3)$	0	~ 0

- Lots of variance in guessing, few solutions are good

Random guessing again



MSE < 0.25



Simple example

- $y = f(w, x)$
- $= w_0 + w_1x_1 + w_2x_2$


x	t	$y=f(x)$
$\sim(3,2)$	1	~ 1
$\sim(-2,-3)$	0	~ 0

- Doesn't scale well, suppose 10 possible values in each dimension:
 - 100 in 2D
 - 1000 in 3D
 - 1,000,000 in 6D...

$w_1=-0.60, w_2=-0.80, b=0.29$

$w_1=-0.91, w_2=-0.44, b=0.34$

Random guessing again




MSE < 0.005



Function as a graph

- $y = f(w, x) = w_0x_0 + w_1x_1 + w_2x_2$

x	t
$\sim(3,2)$	1
$\sim(-2,-3)$	0






Simple example

- $y = f(w, x) = w_0x_0 + w_1x_1 + w_2x_2$

x	t
$\sim(3,2)$	1
$\sim(-2,-3)$	0

$$w = [-1, 0.5, 0.1]$$






Simple example

- $y = f(w, x) = w_0x_0 + w_1x_1 + w_2x_2$

x	t
$\sim(3,2)$	1
$\sim(-2,-3)$	0

$$w = [-1, 0.5, 0.1]$$






Simple example

- $y = f(w, x) = w_0x_0 + w_1x_1 + w_2x_2$

x	t
$\sim(3,2)$	1
$\sim(-2,-3)$	0

$$w = [-1, 0.5, 0.1]$$



Make y_i look like t_i .
More formally, this
means minimize
 $(y_i - t_i)^2$




Simple example

- $y = f(w, x) = w_0x_0 + w_1x_1 + w_2x_2$

x	t
$\sim(3,2)$	1
$\sim(-2,-3)$	0

$$w = [-1, 0.5, 0.1]$$



$t_i: 1$

How do we improve our guess?




Simple example

- $y = f(w, x) = w_0x_0 + w_1x_1 + w_2x_2$

x	t
$\sim(3,2)$	1
$\sim(-2,-3)$	0

$$w = [-1, 0.5, 0.1]$$



Method 1: Fix everything but one variable and adjust single variable

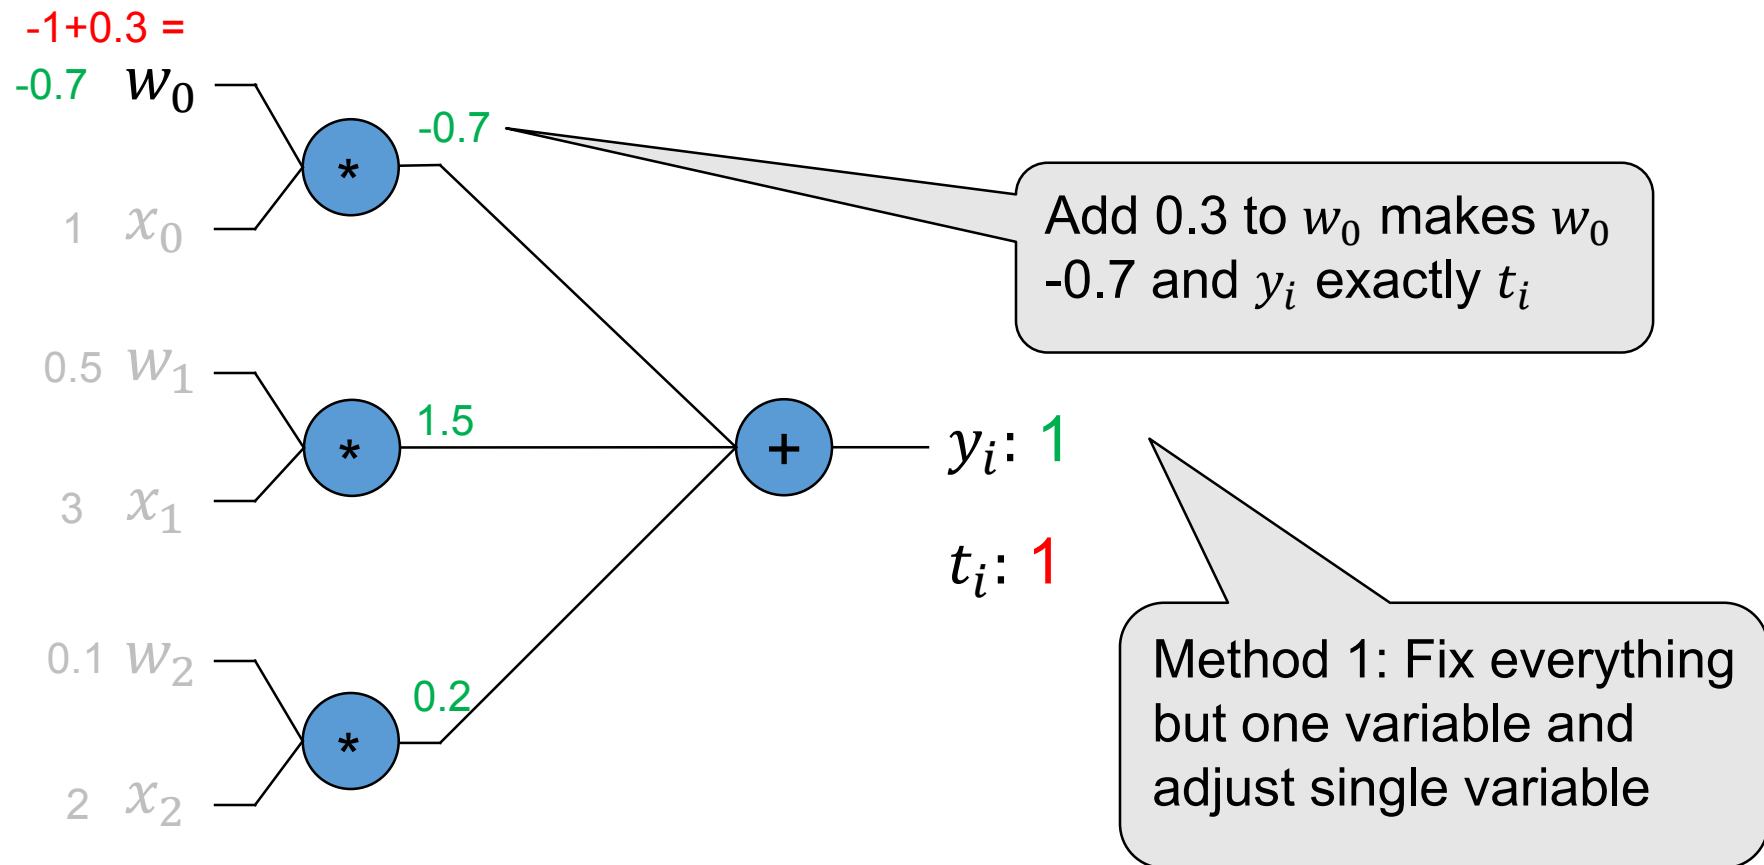


Simple example

- $y = f(w, x) = w_0x_0 + w_1x_1 + w_2x_2$

x	t
$\sim(3,2)$	1
$\sim(-2,-3)$	0

$$w = [-0.7, 0.5, 0.1]$$



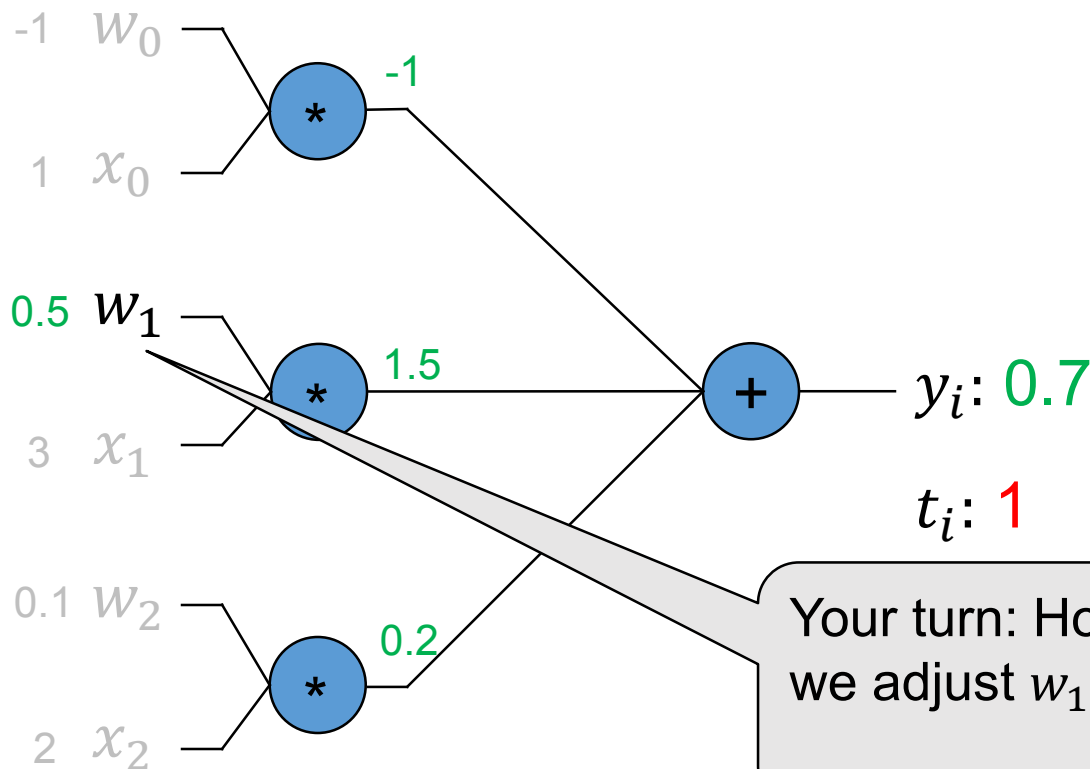


Simple example

- $y = f(w, x) = w_0x_0 + w_1x_1 + w_2x_2$

x	t
$\sim(3,2)$	1
$\sim(-2,-3)$	0

$$w = [-1, 0.5, 0.1]$$



Your turn: How should we adjust w_1 ?

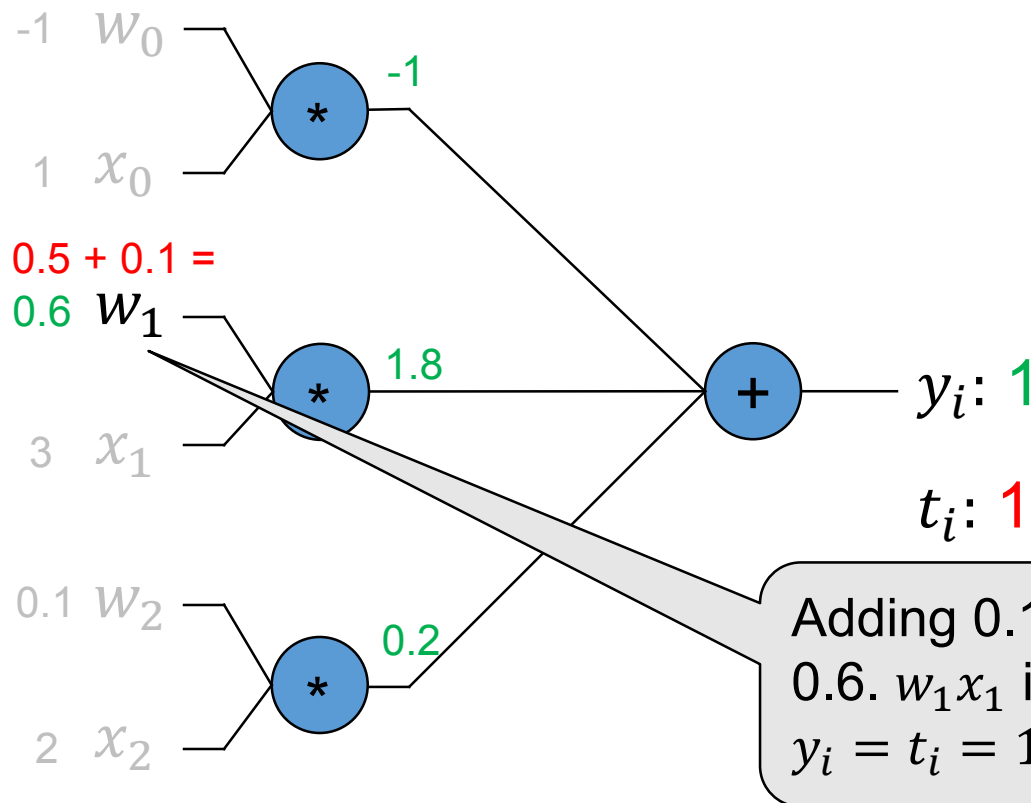


Simple example

- $y = f(w, x) = w_0x_0 + w_1x_1 + w_2x_2$

x	t
$\sim(3,2)$	1
$\sim(-2,-3)$	0

$$w = [-1, 0.6, 0.1]$$



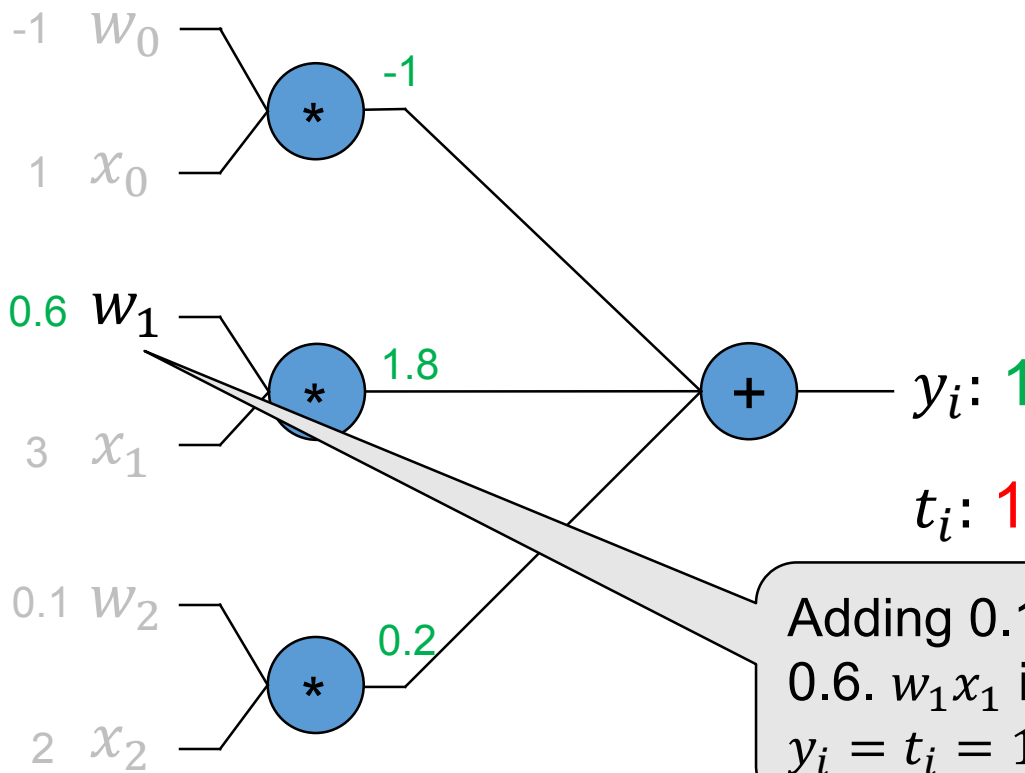


Simple example

- $y = f(w, x) = w_0x_0 + w_1x_1 + w_2x_2$

x	t
$\sim(3,2)$	1
$\sim(-2,-3)$	0

$$w = [-1, 0.6, 0.1]$$




We've basically found an optimization method called *"coordinate descent"*

Adding 0.1 to w_1 makes 0.7. w_1x_1 is now 2.1 and $y_i = t_i = 1$



Coordinate Descent


- Repeat until convergence
 - Choose a coordinate (randomly)
 - Step towards minimum along an axis
 - Update solution vector





But how do we figure out what's “Best”?


- Pros:
 - Simple
- Cons:
 - Slow (updates only one direction at a time)
 - Stepping in one direction can “undo” gains in another





Issues with Coordinate descent

- Updating one dimension at a time is slow
 - We might hit a local minima
 - Evaluating at lots of points is slow
-
- Can we do better?
 - Any ideas?






Issues with Coordinate descent

- Updating one dimension at a time is slow
- We might hit a local minima
- Evaluating at lots of points is slow

- Can we do better?
 - Any ideas?

Use calculus to
minimize!





But how do we figure out what's "Best"?


□ Pros:

Improvement:

Take a step η in the direction of steepest descent (slope)

slope

$$\approx \lim_{h \rightarrow 0} \frac{f(a+h) - f(a)}{h}$$





But how do we figure out what's "Best"?

□ Pros:

Improvement:

Take a step η in the direction of steepest descent (slope)

- Finite difference unbalanced (so use centered difference)

slope

$$\approx \lim_{h \rightarrow 0} \frac{f(a + h) - f(a - h)}{2h}$$

Why?

Finite difference error:

$$O(\Delta h)$$

Centered difference error:

$$O(\Delta h^2)$$



Quick review of derivatives

- $f(x, y) = xy \rightarrow \boxed{\frac{\partial f}{\partial x} = ?, \frac{\partial f}{\partial y} = ?}$
- $\frac{\partial f(x)}{\partial x} = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$
- $f(x + h) = f(x) + h \frac{\partial f(x)}{\partial x}$

Example: $x=5, y=-2 \rightarrow f(x,y)=x*y=?$

$$\boxed{\frac{\partial f}{\partial x} = ?}$$

$$\boxed{\frac{\partial f}{\partial y} = ?}$$

$$\boxed{\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]}$$

Partial derivatives

Gradient



Quick review of derivatives

- $f(x, y) = xy \rightarrow \frac{\partial f}{\partial x} = y, \frac{\partial f}{\partial y} = x$
- $\frac{\partial f(x)}{\partial x} = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$
- $f(x + h) = f(x) + h \frac{\partial f(x)}{\partial x}$

Example: $x=5, y=-2 \rightarrow f(x,y)=x*y=-10$

$$\frac{\partial f}{\partial x} = -2$$

$$\frac{\partial f}{\partial y} = 5$$

Partial derivatives

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

Gradient




Compound expressions

$$f(x, y, z) = (x + y)z$$

$$q = x + y$$

$$f = qz$$



Forward pass:

$$q = x + y \rightarrow q = ?$$

$$f = q * z \rightarrow f = ?$$

$$x=-2, y=5, z=-4$$




Compound expressions

$$f(x, y, z) = (x + y)z$$

$$q = x + y$$

$$f = qz$$



Forward pass:

$$q = x + y \rightarrow q = 3$$

$$f = q * z \rightarrow f = -12$$

$$x=-2, y=5, z=-4$$




Compound expressions

$$f(x, y, z) = (x + y)z = qz$$

$$q = x + y \quad \frac{\partial q}{\partial x} = ?, \quad \frac{\partial q}{\partial y} = ?$$

$$\frac{\partial f}{\partial q} = ?, \quad \frac{\partial f}{\partial z} = ?$$



Forward pass:

$$q = x + y \rightarrow q = 3$$

$$f = q * z \rightarrow f = -12$$

$$x=-2, y=5, z=-4$$




Compound expressions

$$f(x, y, z) = (x + y)z = qz$$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$\frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$



Forward pass:

$$q = x + y \rightarrow q = 3$$
$$f = q * z \rightarrow f = -12$$

$$x=-2, y=5, z=-4$$

Backward pass:

$$dfdz = q \rightarrow 3$$
$$dfdq = z \rightarrow -4$$



Compound expressions


$$f(x, y, z) = (x + y)z = qz$$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$\frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Chain rule:

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial x}$$



Forward pass:

$$q = x + y \rightarrow q = 3$$

$$f = q * z \rightarrow f = -12$$

$$x=-2, y=5, z=-4$$

Backward pass:

$$dfdz = q \rightarrow 3$$

$$dfdq = z \rightarrow -4$$

Back one more level:


$$dfdx = dqdx * dfdq = 1.0 * dfdq \rightarrow -4$$

$$dfdy = dqdy * dfdq = 1.0 * dfdq \rightarrow -4$$




Congratulations!

- We've now discovered the basics of learning weights in a NN take a step along direction of steepest descent (gradient)
 - NN consists of elementary operations such as: $+$, $*$, $/$, e^x
 - Neural network can be expressed in terms of these basic operations
- And we've learned how to optimize one part of these



Schematic of a biological neuron.




Schematic of Rosenblatt's perceptron.

$$f(w, x) = \frac{1}{1 + e^{-s}}$$
$$= \frac{1}{1 + e^{-(w_0 + w_1 x_1 + w_2 x_2)}}$$




Sigmoid example






Sigmoid example






Sigmoid example






Sigmoid example



$f(x) = e^x \rightarrow \frac{\partial f}{\partial x} = ?$	$f(x) = \frac{1}{x} \rightarrow \frac{\partial f}{\partial x} = ?$
$f_a(x) = ax \rightarrow \frac{\partial f}{\partial x} = ?$	$f_c(x) = c + x \rightarrow \frac{\partial f}{\partial x} = ?$




Sigmoid example



$f(x) = e^x \rightarrow \frac{\partial f}{\partial x} = e^x$	$f(x) = \frac{1}{x} \rightarrow \frac{\partial f}{\partial x} = \frac{-1}{x^2}$
$f_a(x) = ax \rightarrow \frac{\partial f}{\partial x} = a$	$f_c(x) = c + x \rightarrow \frac{\partial f}{\partial x} = 1$




Sigmoid example



$f(x) = e^x \rightarrow \frac{\partial f}{\partial x} = e^x$	$f(x) = \frac{1}{x} \rightarrow \frac{\partial f}{\partial x} = \frac{-1}{x^2}$
$f_a(x) = ax \rightarrow \frac{\partial f}{\partial x} = a$	$f_c(x) = c + x \rightarrow \frac{\partial f}{\partial x} = 1$




Sigmoid example



$f(x) = e^x \rightarrow \frac{\partial f}{\partial x} = e^x$	$f(x) = \frac{1}{x} \rightarrow \frac{\partial f}{\partial x} = \frac{-1}{x^2}$
$f_a(x) = ax \rightarrow \frac{\partial f}{\partial x} = a$	$f_c(x) = c + x \rightarrow \frac{\partial f}{\partial x} = 1$



Sigmoid example



$$f(x) = e^x \rightarrow \frac{\partial f}{\partial x} = e^x$$


$$f(x) = \frac{1}{x} \rightarrow \frac{\partial f}{\partial x} = \frac{-1}{x^2}$$

$$f_a(x) = ax \rightarrow \frac{\partial f}{\partial x} = a$$

$$f_c(x) = c + x \rightarrow \frac{\partial f}{\partial x} = 1$$




Sigmoid example



$f(x) = e^x \rightarrow \frac{\partial f}{\partial x} = e^x$	$f(x) = \frac{1}{x} \rightarrow \frac{\partial f}{\partial x} = \frac{-1}{x^2}$
$f_a(x) = ax \rightarrow \frac{\partial f}{\partial x} = a$	$f_c(x) = c + x \rightarrow \frac{\partial f}{\partial x} = 1$



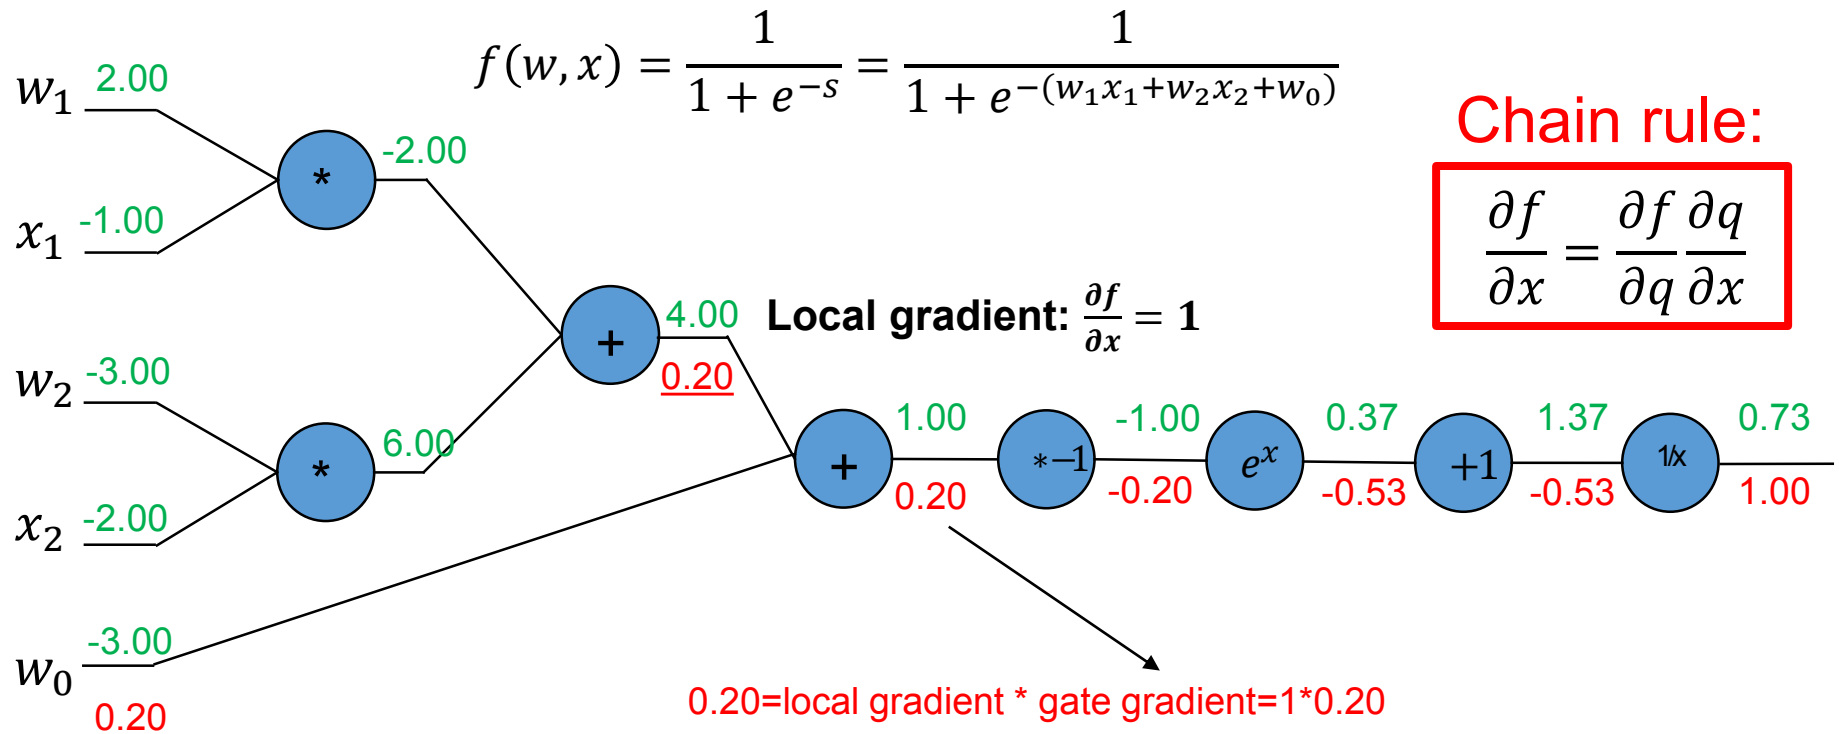
Sigmoid example



$f(x) = e^x \rightarrow \frac{\partial f}{\partial x} = e^x$	$f(x) = \frac{1}{x} \rightarrow \frac{\partial f}{\partial x} = \frac{-1}{x^2}$
$f_a(x) = ax \rightarrow \frac{\partial f}{\partial x} = a$	$f_c(x) = c + x \rightarrow \frac{\partial f}{\partial x} = 1$




Sigmoid example



$f(x) = e^x \rightarrow \frac{\partial f}{\partial x} = e^x$	$f(x) = \frac{1}{x} \rightarrow \frac{\partial f}{\partial x} = \frac{-1}{x^2}$
$f_a(x) = ax \rightarrow \frac{\partial f}{\partial x} = a$	$f_c(x) = c + x \rightarrow \frac{\partial f}{\partial x} = 1$



Sigmoid example



$$f(x) = e^x \rightarrow \frac{\partial f}{\partial x} = e^x$$


$$f(x) = \frac{1}{x} \rightarrow \frac{\partial f}{\partial x} = \frac{-1}{x^2}$$

$$f_a(x) = ax \rightarrow \frac{\partial f}{\partial x} = a$$

$$f_c(x) = c + x \rightarrow \frac{\partial f}{\partial x} = 1$$



Sigmoid example



$$f(x) = e^x \rightarrow \frac{\partial f}{\partial x} = e^x$$


$$f(x) = \frac{1}{x} \rightarrow \frac{\partial f}{\partial x} = \frac{-1}{x^2}$$

$$f_a(x) = ax \rightarrow \frac{\partial f}{\partial x} = a$$

$$f_c(x) = c + x \rightarrow \frac{\partial f}{\partial x} = 1$$



Sigmoid example



This part is just our
simple linear classifier:

$$w_1x_1 + w_2x_2 + w_0$$



$$f(x) = e^x \rightarrow \frac{\partial f}{\partial x} = e^x$$

$$f(x) = \frac{1}{x} \rightarrow \frac{\partial f}{\partial x} = \frac{-1}{x^2}$$

$$f_a(x) = ax \rightarrow \frac{\partial f}{\partial x} = a$$

$$f_c(x) = c + x \rightarrow \frac{\partial f}{\partial x} = 1$$




Sigmoid example

$$f(w, x) = \frac{1}{1 + e^{-s}} = \frac{1}{1 + e^{-(w_1x_1 + w_2x_2 + w_0)}}$$

This part is just a
sigmoid activation:

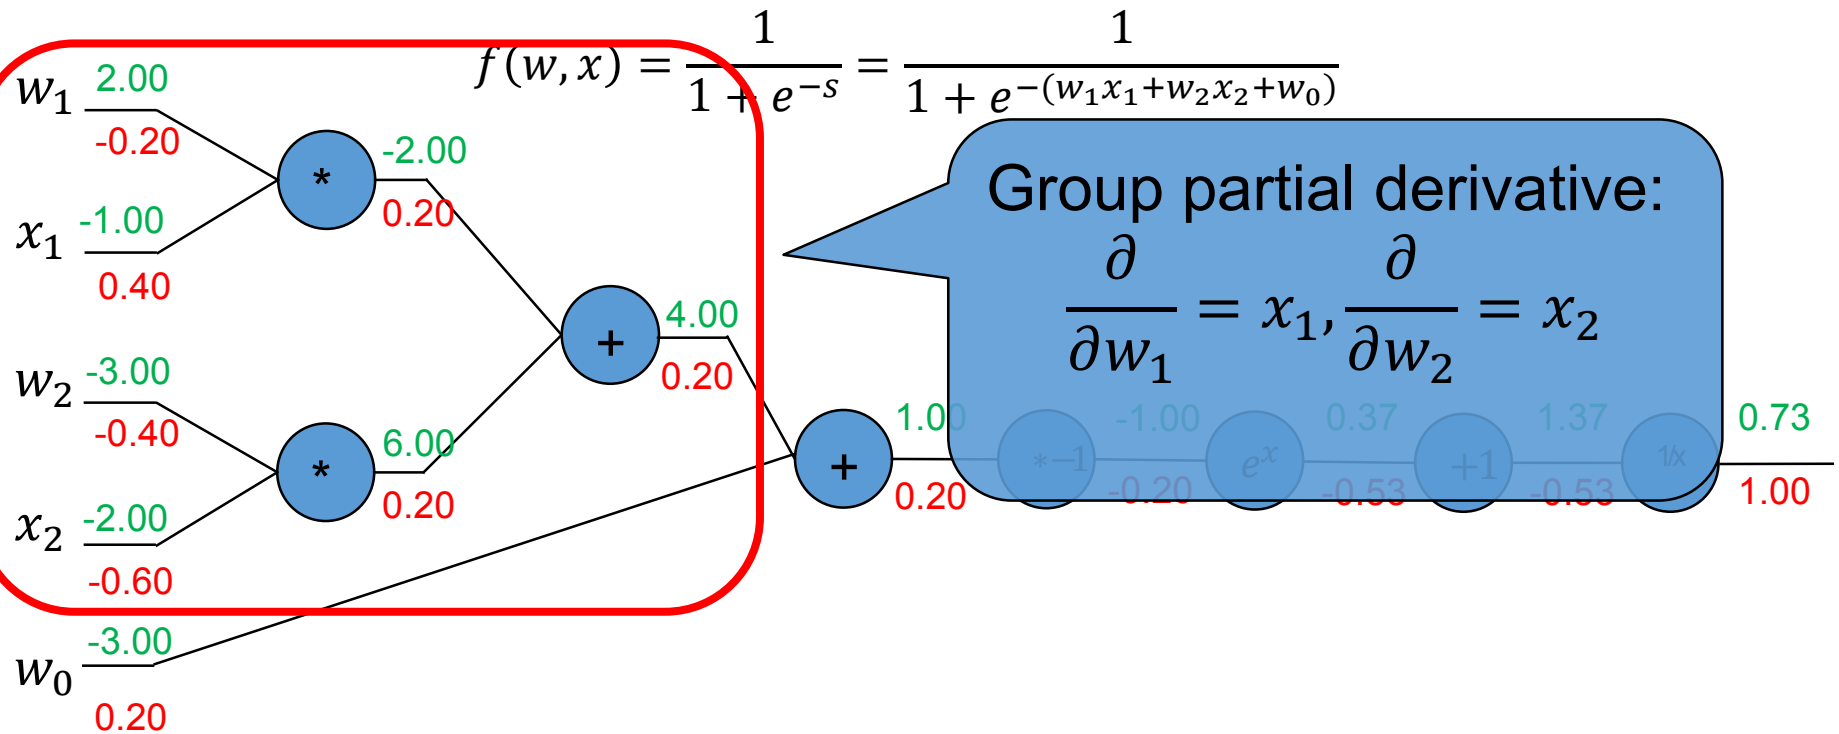
$$\frac{1}{1 + e^{-s}}$$



$f(x) = e^x \rightarrow \frac{\partial f}{\partial x} = e^x$	$f(x) = \frac{1}{x} \rightarrow \frac{\partial f}{\partial x} = \frac{-1}{x^2}$
$f_a(x) = ax \rightarrow \frac{\partial f}{\partial x} = a$	$f_c(x) = c + x \rightarrow \frac{\partial f}{\partial x} = 1$



Sigmoid example



$f(x) = e^x \rightarrow \frac{\partial f}{\partial x} = e^x$	$f(x) = \frac{1}{x} \rightarrow \frac{\partial f}{\partial x} = \frac{-1}{x^2}$
$f_a(x) = ax \rightarrow \frac{\partial f}{\partial x} = a$	$f_c(x) = c + x \rightarrow \frac{\partial f}{\partial x} = 1$



Sigmoid example

$$f(w, x) = \frac{1}{1 + e^{-s}} = \frac{1}{1 + e^{-(w_1 x_1 + w_2 x_2 + w_0)}}$$

Block partial derivative:

$$\frac{\partial}{\partial s} \frac{1}{1 + e^{-s}} = -1(1 + e^{-s})^{-2} \frac{\delta}{\delta s} (1 + e^{-s}) = -1(1 + e^{-s})^{-2} (-1e^{-s}) = \frac{e^{-s}}{(1 + e^{-s})^2} = \frac{1}{1 + e^{-s}} \left(\frac{1 + e^{-s}}{1 + e^{-s}} - \frac{1}{1 + e^{-s}} \right) = \sigma(s)(1 - \sigma(s))$$

w_1 2.00
-0.20
-2.00
-1.00
0.40
3.00
-0.40
-0.60
-3.00
0.20

x_2 -2.00
-0.60
-3.00
0.20

w_0 -3.00
0.20


0.37
1.37
0.53
0.73
1.00

$1/x$

$f(x) = e^x \rightarrow \frac{\partial f}{\partial x} = e^x$	$f(x) = \frac{1}{x} \rightarrow \frac{\partial f}{\partial x} = \frac{-1}{x^2}$
$f_a(x) = ax \rightarrow \frac{\partial f}{\partial x} = a$	$f_c(x) = c + x \rightarrow \frac{\partial f}{\partial x} = 1$



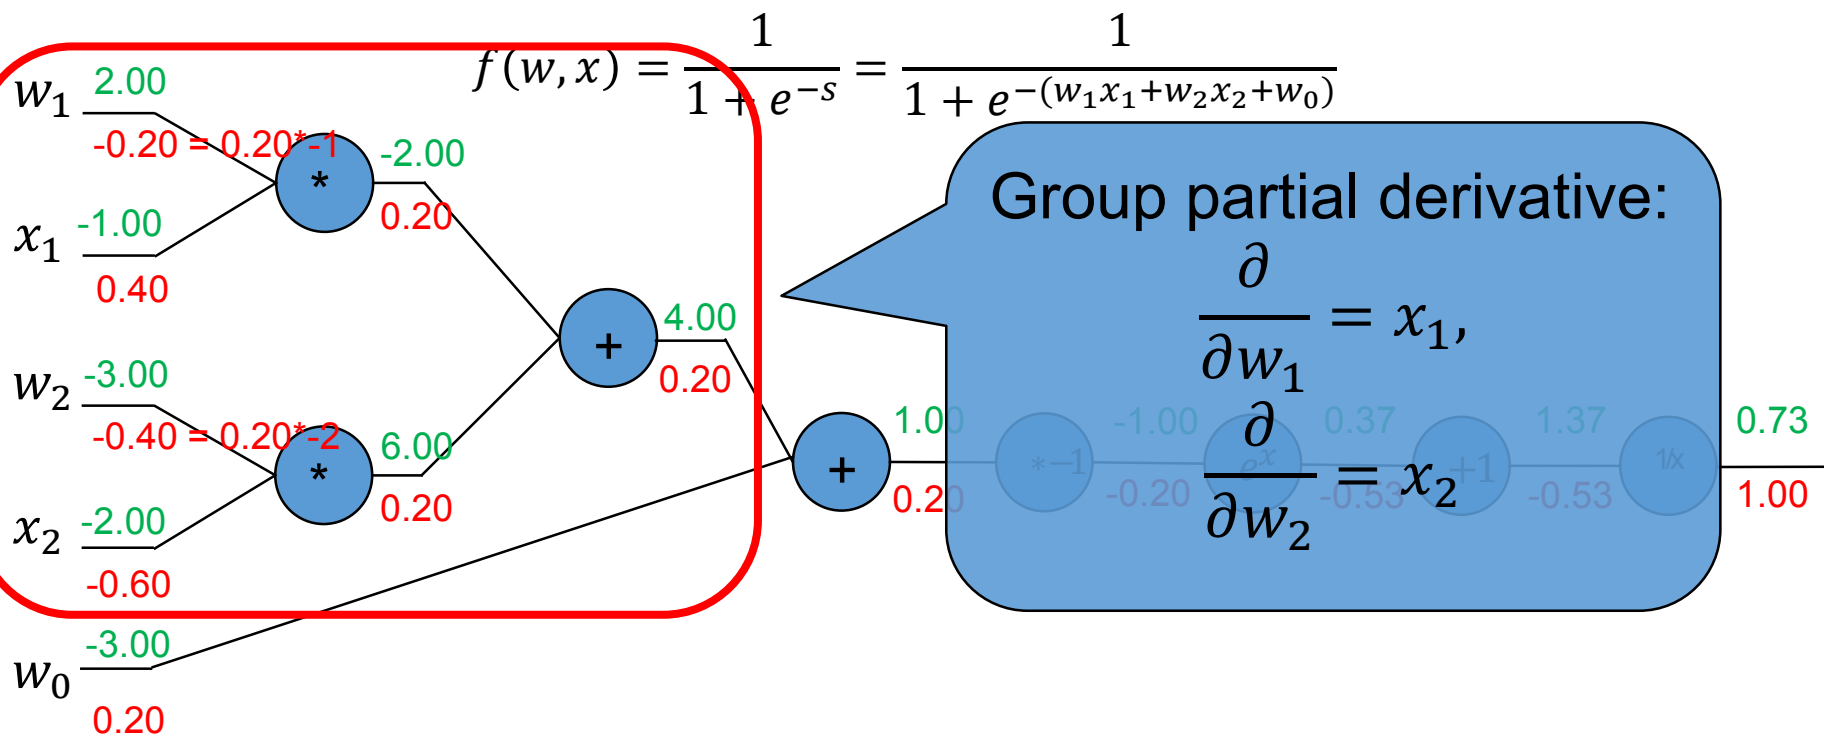
Sigmoid example



$f(x) = e^x \rightarrow \frac{\partial f}{\partial x} = e^x$	$f(x) = \frac{1}{x} \rightarrow \frac{\partial f}{\partial x} = \frac{-1}{x^2}$
$f_a(x) = ax \rightarrow \frac{\partial f}{\partial x} = a$	$f_c(x) = c + x \rightarrow \frac{\partial f}{\partial x} = 1$



Sigmoid example




$f(x) = e^x \rightarrow \frac{\partial f}{\partial x} = e^x$	$f(x) = \frac{1}{x} \rightarrow \frac{\partial f}{\partial x} = \frac{-1}{x^2}$
$f_a(x) = ax \rightarrow \frac{\partial f}{\partial x} = a$	$f_c(x) = c + x \rightarrow \frac{\partial f}{\partial x} = 1$



Sigmoid example

$$f(w, x) = \frac{1}{1 + e^{-s}} = \frac{1}{1 + e^{-(w_1x_1 + w_2x_2 + w_0)}}$$



Gradient on w_1 :
 $\sigma(1 - \sigma) * x_1$
Gradient on w_2 :
 $\sigma(1 - \sigma) * x_2$

$f(x) = e^x \rightarrow \frac{\partial f}{\partial x} = e^x$	$f(x) = \frac{1}{x} \rightarrow \frac{\partial f}{\partial x} = \frac{-1}{x^2}$
$f_a(x) = ax \rightarrow \frac{\partial f}{\partial x} = a$	$f_c(x) = c + x \rightarrow \frac{\partial f}{\partial x} = 1$



- Gradient on MSE:
- $MSE = (y_i - t_i)^2$
- $\frac{\partial}{\partial y} MSE = 2 * (y_i - t_i)$
- MSE gradient * Sigmoid gradient:
- Overall gradient=MSE gradient * Sigmoid gradient:
- $2(\sigma - t) * \sigma(1 - \sigma) * x_1$
MSE Sigmoid Dot product



Now that we know the gradient...

□ Pros:

Improvement:

Take a step η in the direction of steepest descent as determined by the analytical gradient

Use numeric gradient to double-check analytical gradient

slope

$$\approx \lim_{h \rightarrow 0} \frac{f(a+h) - f(a-h)}{2h}$$

Why?

Finite difference error:

$$O(\Delta h)$$


Centered difference error:

$$O(\Delta h^2)$$




Sigmoid example


- Our long example was just a single neuron
- This neuron can be trained using backpropagation




Schematic of a biological neuron.



Schematic of Rosenblatt's perceptron.



$$f(w, x) = \frac{1}{1 + e^{-s}} = \frac{1}{1 + e^{-(w_1x_1 + w_2x_2 + w_0)}}$$

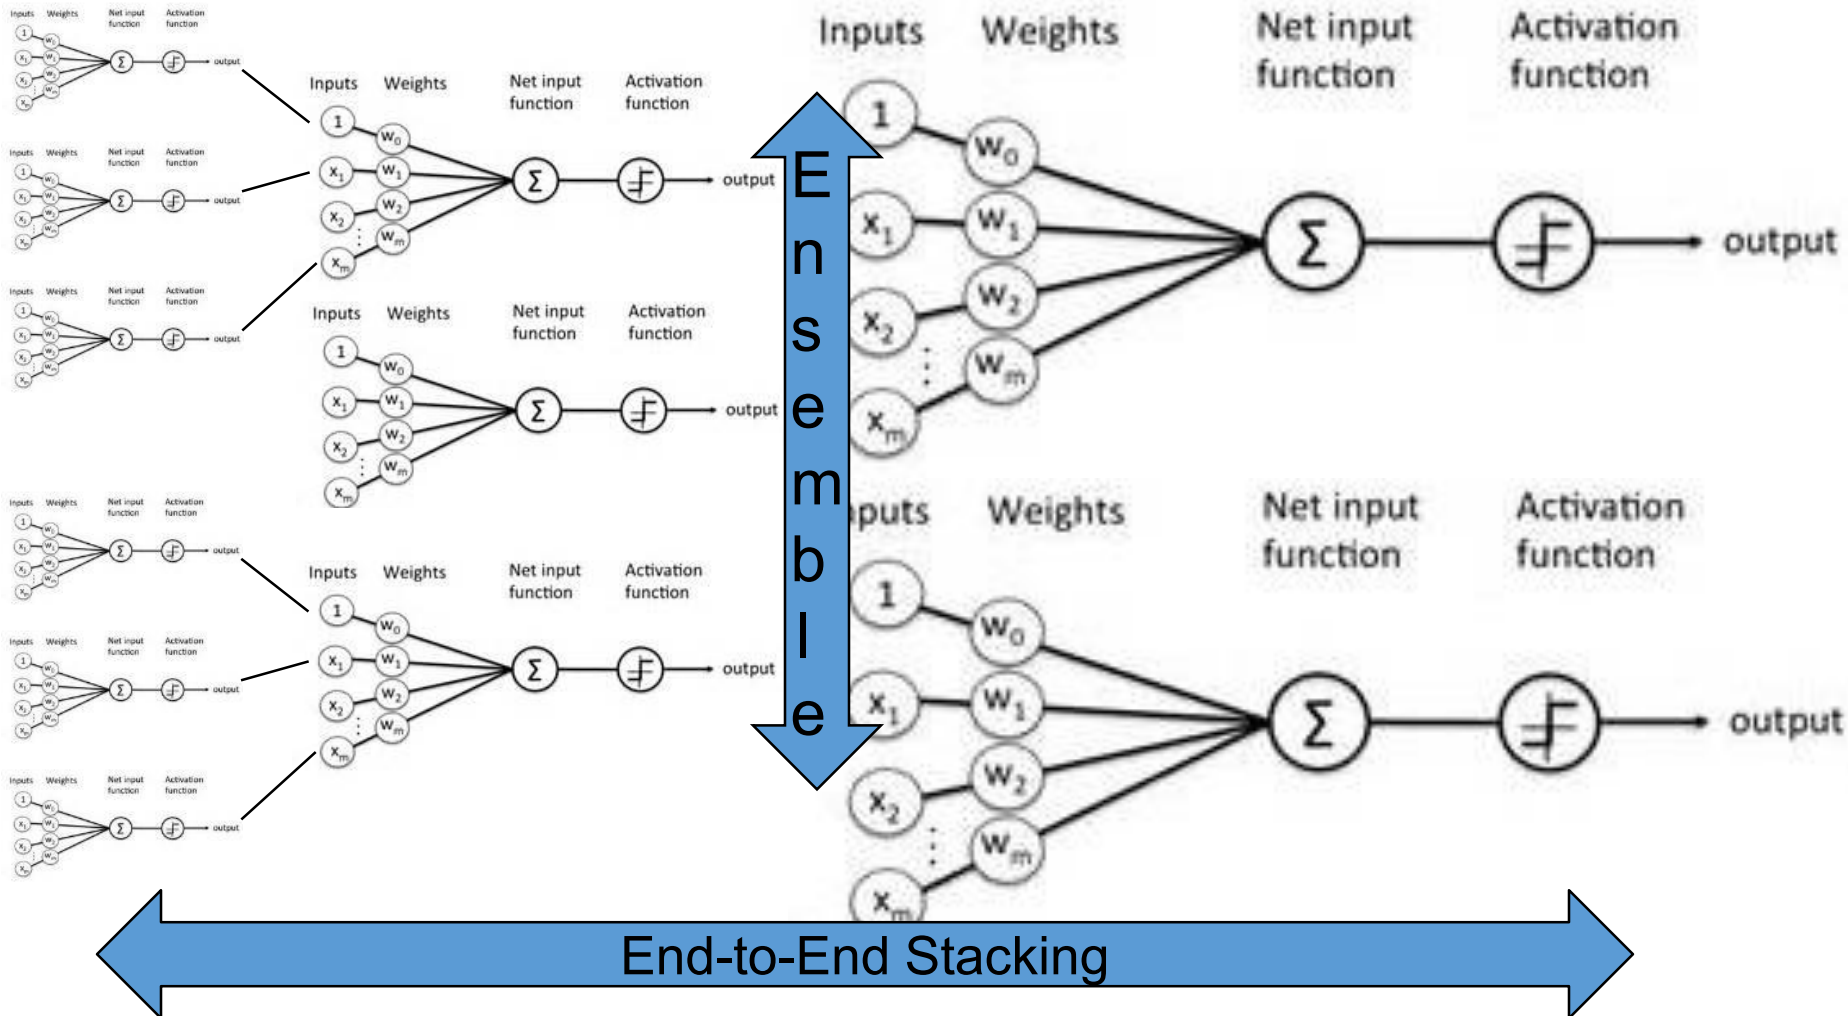


$$w_0 = -3.00$$
$$0.20$$




NORTHWESTERN
UNIVERSITY

“Layered” Logistic Regression






Feedforward networks




Goal: Learn weights such that outputs agree with training data \rightarrow Minimize $\sum_i (Y_i - O_i)^2$



Feedforward networks




0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19
20	21	22	23	24



Training
Value Y_i

Output O_i



Ground truth:
Make prediction
match this

Goal: Learn weights such that outputs agree with training data \rightarrow Minimize $\sum_i (Y_i - O_i)^2$

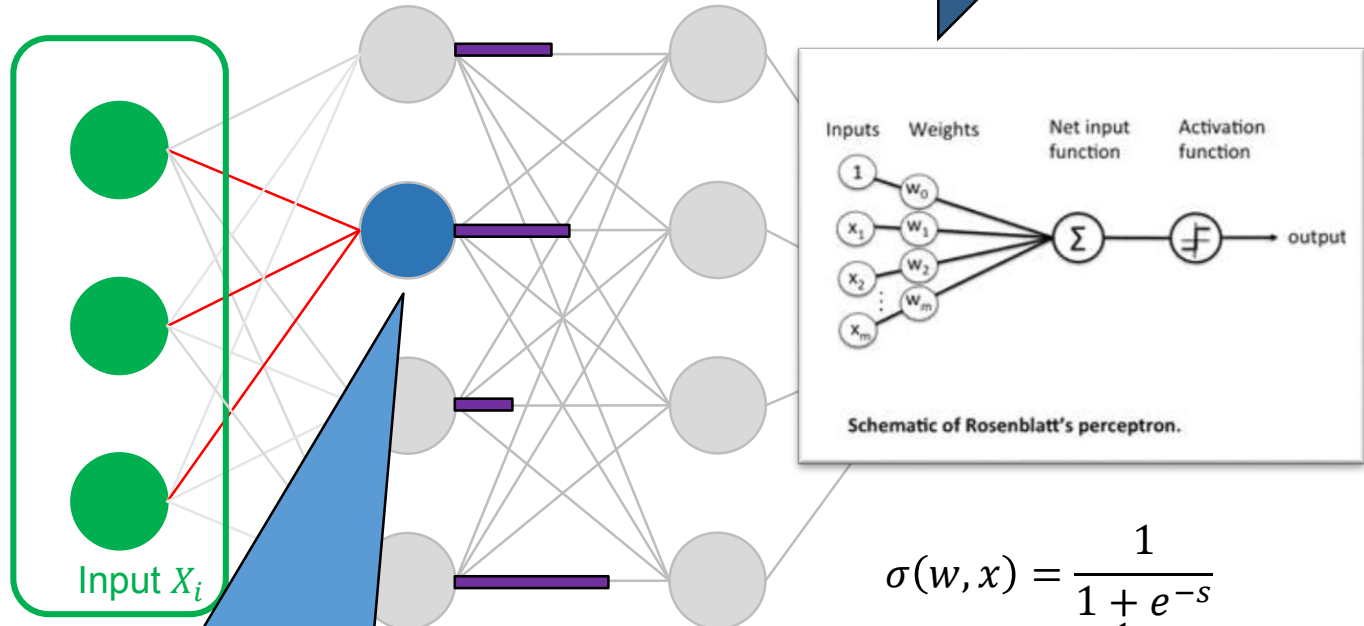


NORTHWESTERN
UNIVERSITY

Forward propagation

Forward propagation


0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19
20	21	22	23	24



$$\sigma(w, x) = \frac{1}{1 + e^{-s}}$$
$$= \frac{1}{1 + e^{-(w_0 + w_1 x_1 + w_2 x_2)}}$$

Each activation is simply a logistic regression.

Bar length = level of activation






NORTHWESTERN
UNIVERSITY

Forward propagation

Forward propagation

0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19
20	21	22	23	24



Each activation is simply a stacked logistic regression.


Bar length = level of activation



Forward propagation


Forward propagation

0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19
20	21	22	23	24






Forward propagation





Backpropagation




We must adjust the predictions using error derivatives → Backpropagation

Goal: Learn weights such that outputs agree with training data → Minimize $\sum_i (Y_i - O_i)^2$



NORTHWESTERN
UNIVERSITY

Backpropagation



Recall the following activations in the forward pass



NORTHWESTERN
UNIVERSITY

Backpropagation




Figure out how to adjust input(**green**) weights(**red**) to match target activations(**purple**)



Backpropagation

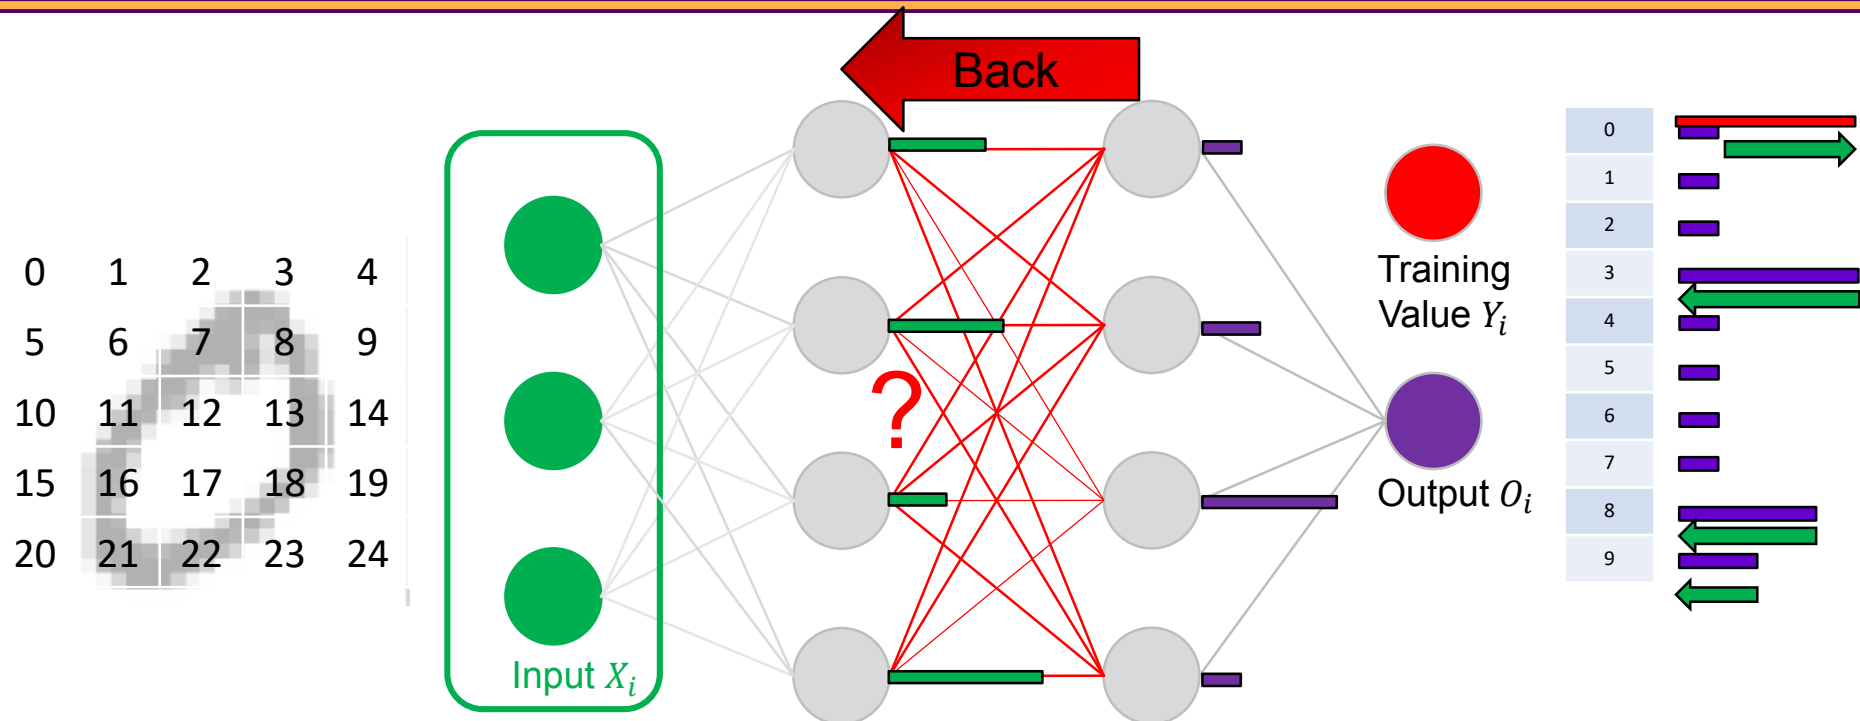
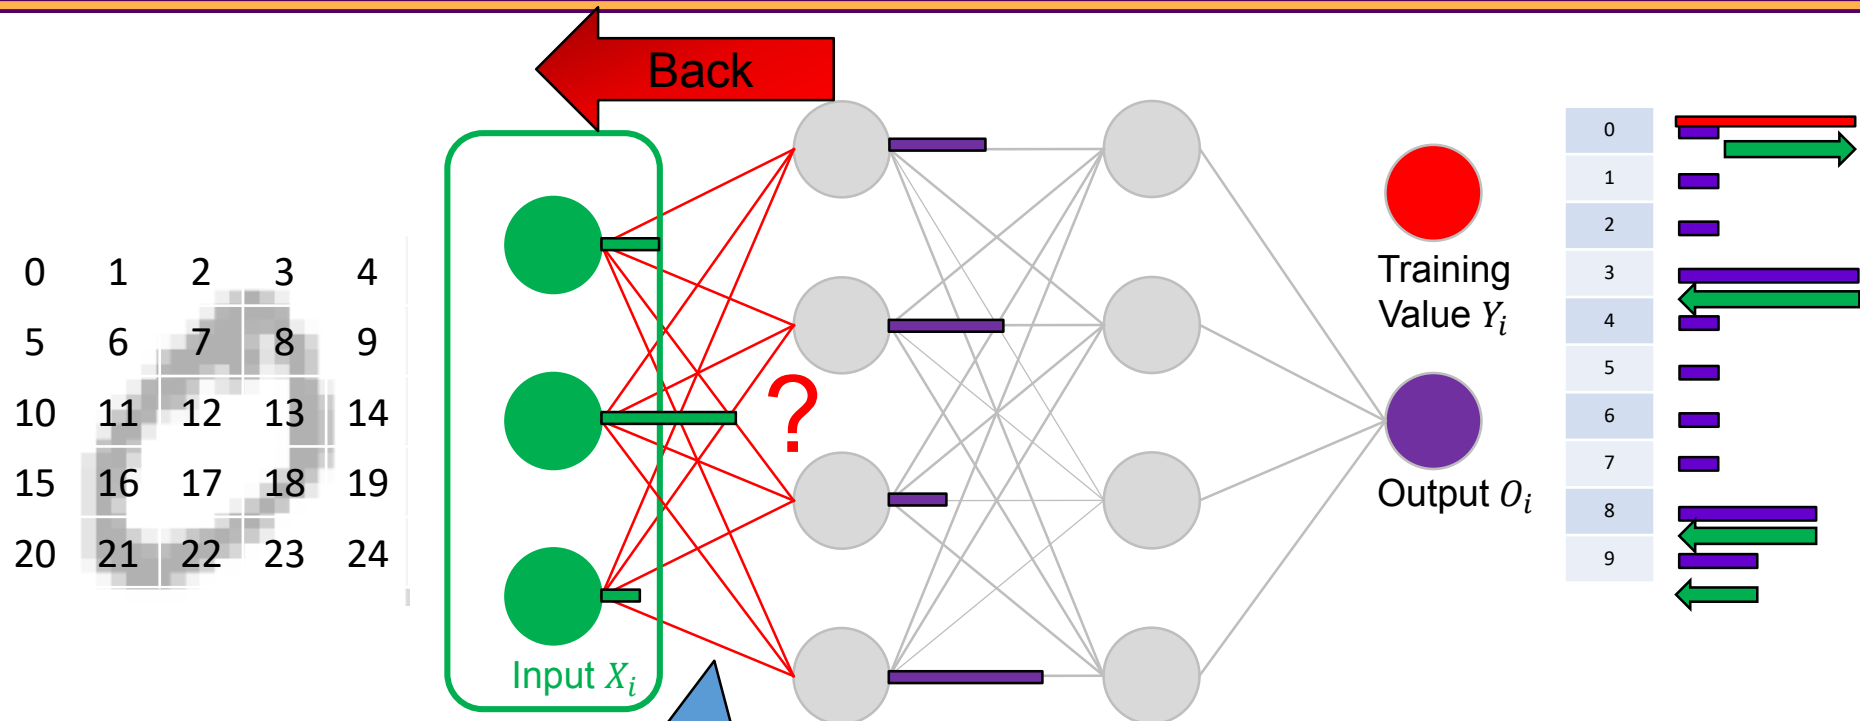


Figure out how to adjust
input(**green**) weights(**red**) to
match target activations(**purple**)



Backpropagation






NORTHWESTERN
UNIVERSITY


MSIA 490-30: Deep Learning. Spring 2017.
Instructor: Dr. Ellick Chan. TA: Mark Harmon.

Feedforward networks






Feedforward networks



Each input neuron carries a signal
neurons in input layer
= # inputs
= #pixels = 25

Recall: vectorizing inputs


0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19
20	21	22	23	24



Samples are column
vectors



Feedforward networks



$$\sigma \left(\begin{matrix} w_{1,1} & w_{1,2} & w_{1,3} \\ x_1 \\ x_2 \\ x_3 \end{matrix} + w_0 \right) = a_1$$


- Signals are propagated to the second layer
- Each neuron in the second layer is a logistic regression

w_1
...
...
...
w_n

Weights are row vectors



Feedforward networks




$$\sigma \left(\begin{matrix} w_{1,1} & w_{1,2} & w_{1,3} \\ w_{1,4} & w_{1,5} & w_{1,6} \\ w_{1,7} & w_{1,8} & w_{1,9} \\ w_{1,10} & w_{1,11} & w_{1,12} \end{matrix} \begin{matrix} x_1 \\ x_2 \\ x_3 \end{matrix} + w_0 \right) = \begin{matrix} a_1 \\ a_2 \\ a_3 \\ a_4 \end{matrix}$$

- Signals are propagated to the second layer \rightarrow Vectorize
- Each neuron in the second layer is a logistic regression

```
def sigm(s):  
    return 1.0/(1 + np.e**-s)  
A1 = sigm(np.dot(W1, X))
```



Feedforward networks




$$\sigma \left(W_2 \begin{bmatrix} h_1 \\ h_2 \\ h_3 \\ h_4 \end{bmatrix} + w_0 \right) = \begin{bmatrix} a_1 \\ a_2 \\ a_3 \\ a_4 \end{bmatrix}$$

- Signals are propagated to the third layer
- Each neuron in the third layer is a logistic regression

```
def sigm(s):  
    return 1.0/(1 + np.e**(-s))  
A1 = sigm(np.dot(W1, X))  
A2 = sigm(np.dot(W2, A1))
```



Feedforward networks



$$\sigma \left(\begin{bmatrix} & W_3 & \end{bmatrix} \begin{bmatrix} h_1 \\ h_2 \\ h_3 \\ h_4 \end{bmatrix} + w_0 \right) = \text{Out}$$

- Signals are propagated to the final layer
- The output is a logistic regression

```
def sigm(s):  
    return 1.0/(1 + np.e**-s)  
A1 = sigm(np.dot(W1, X))  
A2 = sigm(np.dot(W2, A1))  
Out = sigm(np.dot(W3, A2))
```



NORTHWESTERN
UNIVERSITY

MSIA 490-30: Deep Learning. Spring 2017.
Instructor: Dr. Ellick Chan. TA: Mark Harmon.

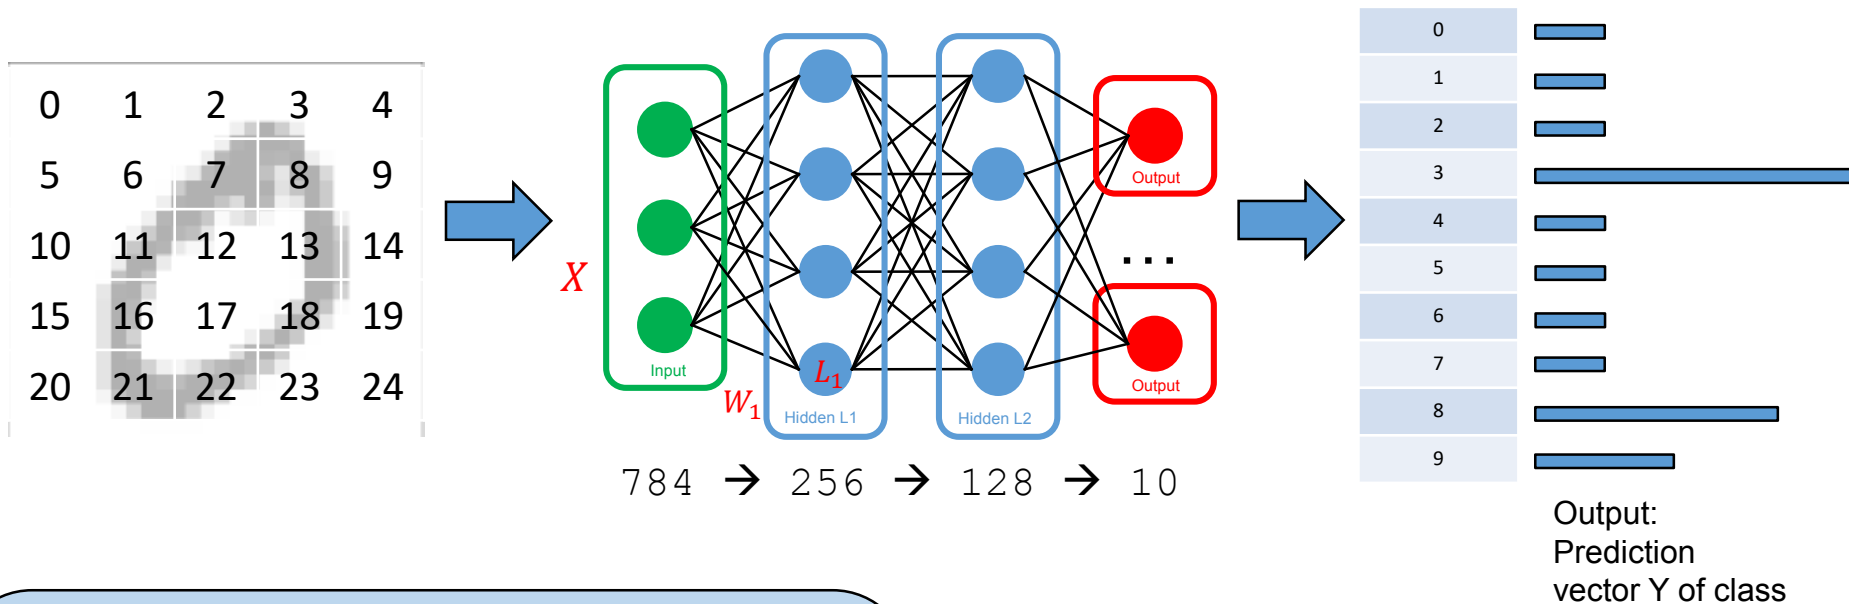
Making things fast

VECTORIZING FEEDFORWARD NETS



NORTHWESTERN
UNIVERSITY

Vectorized forward pass



Setup code:

```
# 784  $\rightarrow$  256  $\rightarrow$  128  $\rightarrow$  10  
X = training_data # 784 x 60000  
T = target_classes # 10 x 60000
```

```
W1 = 2*rand(784, 256).T - 1  
W2 = 2*rand(256, 128).T - 1  
W3 = 2*rand(128, 10).T - 1
```

A.T means A
transpose: $A_{ij}^T = A_{ji}$

Don't confuse this
with target T




NORTHWESTERN
UNIVERSITY


Vectorized forward pass

Code:


```
L1 = sigmoid(W1.dot(X))
```



σ




Weights: 256 x 784
(256 kernels operating on
784 dimensional images)



Input: 60,000 28x28 images
= 784 pixel vectors
(784 x 60,000)

=



Activations: 256 x 60,000
(256 dimension activations
For 60,000 images)




NORTHWESTERN
UNIVERSITY

Vectorized forward pass

Code:


```
L1 = sigmoid(W1.dot(X))  
L2 = sigmoid(W2.dot(L1))
```



σ




Weights: 128 x 256
(128 kernels operating on
256 dimensional activations)



Input: 60,000 x 256
dimensional activations
(256 x 60,000)

$=$



Activations: 128 x 60,000
(128 dimension activations
for 60,000 samples)




NORTHWESTERN
UNIVERSITY


Vectorized forward pass

Code:


```
L1 = sigmoid(W1.dot(X))  
L2 = sigmoid(W2.dot(L1))  
L3 = sigmoid(W3.dot(L2))
```




$\hat{=}$



σ




Weights: 10 x 128
(10 kernels operating on
128 dimensional activations)



Input: 60,000 x 128
dimensional activations
(128 x 60,000)

$=$



Activations: 10 x 60,000
(10 dimensional classification
confidences for 60,000 samples)




NORTHWESTERN
UNIVERSITY

MSIA 490-30: Deep Learning. Spring 2017.
Instructor: Dr. Ellick Chan. TA: Mark Harmon.

Vectorized backward pass

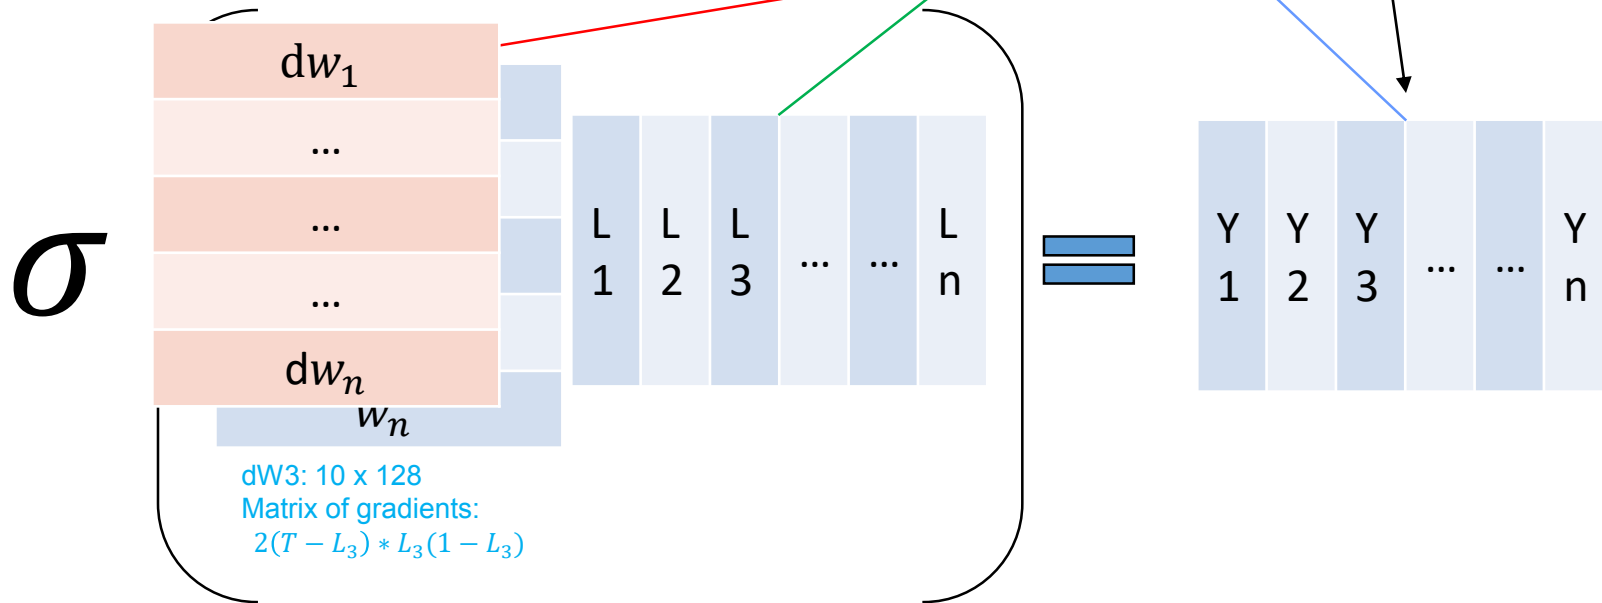
Code:

```
L1 = sigmoid(W1.dot(X))
L2 = sigmoid(W2.dot(L1))
L3 = sigmoid(W3.dot(L2))
dW3 = (L3 - T) * L3*(1 - L3)
```



?

Target
T



dW3: 10 x 128
Matrix of gradients:
 $2(T - L_3) * L_3(1 - L_3)$




NORTHWESTERN
UNIVERSITY

MSIA 490-30: Deep Learning. Spring 2017.
Instructor: Dr. Ellick Chan. TA: Mark Harmon.



Vectorized backward pass

Code:

```
L1 = sigmoid(W1.dot(X))
L2 = sigmoid(W2.dot(L1))
L3 = sigmoid(W3.dot(L2))
dW3 = (L3 - T) * L3*(1 - L3)
dW2 = W3.T.dot(dW3) * (L2*(1-L2))
```



\neq



Need to transpose W_3 in the backward pass to accommodate gradients flowing backward




NORTHWESTERN
UNIVERSITY

MSIA 490-30: Deep Learning. Spring 2017.
Instructor: Dr. Ellick Chan. TA: Mark Harmon.


Vectorized backward pass

Code:


```
L1 = sigmoid(W1.dot(X))
L2 = sigmoid(W2.dot(L1))
L3 = sigmoid(W3.dot(L2))
dW3 = (L3 - T) * L3*(1 - L3)
dW2 = W3.T.dot(dW3) * (L2*(1-L2))
dW1 = W2.T.dot(dW2) * (L1*(1-L1))
```



\neq



σ



dW1: 256 x 784
Matrix of gradients:
 $W_2^T dW_2 * L_1(1 - L_1)$


Need to transpose W_2 in the backward pass to accommodate gradients flowing backward




Vectorized weight updates

Code:

```
L1 = sigmoid(W1.dot(X))
L2 = sigmoid(W2.dot(L1))
L3 = sigmoid(W3.dot(L2))
dW3 = (L3 - T) * L3*(1 - L3)
dW2 = W3.T.dot(dW3) * (L2*(1-L2))
dW1 = W2.T.dot(dW2) * (L1*(1-L1))
W3 -= lr*np.dot(dW3, L2.T)
W2 -= lr*np.dot(dW2, L1.T)
W1 -= lr*np.dot(dW1, X.T)
```



?




$$2(\sigma - t) * \sigma(1 - \sigma) * x_1$$

MSE

Sigmoid

Dot product







Vectorized weight updates

Code:

```
L1 = sigmoid(W1.dot(X))
L2 = sigmoid(W2.dot(L1))
L3 = sigmoid(W3.dot(L2))
dW3 = (L3 - T) * L3 * (1 - L3)
dW2 = W3.T.dot(dW3) * (L2 * (1 - L2))
dW1 = W2.T.dot(dW2) * (L1 * (1 - L1))
W3 -= lr * np.dot(dW3, L2.T)
W2 -= lr * np.dot(dW2, L1.T)
W1 -= lr * np.dot(dW1, X.T)
```



?



Target
T


$$2(\sigma - t) * \sigma(1 - \sigma) * x_1$$

MSE

Sigmoid

Dot prod

NB: Factor of 2 in MSE is absorbed in learning rate





Vectorized weight updates

Code:

```
# Setup
# 784 → 256 → 128 → 10
X = training_data # 784 x 60000
T = target_classes # 10 x 60000


W1 = 2*rand(784, 256).T - 1
W2 = 2*rand(256, 128).T - 1
W3 = 2*rand(128, 10).T - 1

lr = 1e-5
def sigmoid(x): return 1.0/(1.0 + np.e**(-x))
for i in range(5000):
    # Forward pass
    L1 = sigmoid(W1.dot(X))
    L2 = sigmoid(W2.dot(L1))
    L3 = sigmoid(W3.dot(L2))

    # Backward pass
    dW3 = (L3 - T) * L3*(1 - L3)
    dW2 = W3.T.dot(dW3) * (L2*(1-L2))
    dW1 = W2.T.dot(dW2) * (L1*(1-L1))

    # Update
    W3 -= lr*np.dot(dW3, L2.T)
    W2 -= lr*np.dot(dW2, L1.T)
    W1 -= lr*np.dot(dW1, X.T)

    print("[%04d] MSE Loss: %0.3f" % (i, np.sum((L3 - T)**2)/len(T.T)))
```





Assignment schedule

- Assignment 1: Due April 26th
 - Basic NN concepts
 - Optimization
- Assignment 2 (Group): Due May 18th
 - Computer Vision
 - Preparation for project
- Mini Quiz 1: May 8th: Basic NN, optimization and generalization multiple choice
- Mini Quiz 2: June 9th (2nd half): Generative, recurrent and applications multiple choice

April 2017

Su	Mo	Tu	We	Th	Fr	Sa
26	27	28	29	30	31	1
2	3	4	5	6	7	8
9	10	11	12	13	14	15
16	17	18	19	20	21	22
23	24	25	26	27	28	29
30	1	2	3	4	5	6

May 2017

Su	Mo	Tu	We	Th	Fr	Sa
30	1	2	3	4	5	6
7	8	9	10	11	12	13
14	15	16	17	18	19	20
21	22	23	24	25	26	27
28	29	30	31	1	2	3
4	5	6	7	8	9	10

June 2017

Su	Mo	Tu	We	Th	Fr	Sa
28	29	30	31	1	2	3
4	5	6	7	8	9	10
11	12	13	14	15	16	17
18	19	20	21	22	23	24
25	26	27	28	29	30	1
2	3	4	5	6	7	8



NORTHWESTERN
UNIVERSITY

MSIA 490-30: Deep Learning. Spring 2017.
Instructor: Dr. Ellick Chan. TA: Mark Harmon.

Projects

- Project group size
- Poster invitation
- Project proposal



NORTHWESTERN
UNIVERSITY

MSIA 490-30: Deep Learning. Spring 2017.
Instructor: Dr. Ellick Chan. TA: Mark Harmon.


Homework

- HW1



Next time

- Next time:
 - Advanced optimization techniques



- Reading:
<http://neuralnetworksanddeeplearning.com/chap3.html>



New stuff

- Genetic algorithms to design NN:
https://en.wikipedia.org/wiki/Neuroevolution_of_augmenting_topologies
- Overview of various NN:
<https://culurciello.github.io/tech/2016/06/04/nets.html>